

notebook

June 12, 2025

1 1. Exploratory Data Analysis (EDA)

1.1 1.1. Setup

First, let's import the necessary libraries for data manipulation and visualization and load the dataset.

```
[1]: import pandas as pd
import numpy as np
import matplotlib.pyplot as plt
import seaborn as sns
```

1.2 1.2. Initial Data Inspection

Let's get a first look at the data's structure, content, and statistical summary. For that define the paths:

```
[2]: BANK_PATH = "data/bank/bank.csv"
BANK_FULL_PATH = "data/bank/bank-full.csv"
BANK_ADDITIONAL_PATH = "data/bank-additional/bank-additional.csv"
BANK_ADDITIONAL_FULL_PATH = "data/bank-additional/bank-additional-full.csv"
```

```
[3]: bank_df = pd.read_csv(BANK_PATH, sep=";")
bank_full_df = pd.read_csv(BANK_FULL_PATH, sep=";")
bank_additional_df = pd.read_csv(BANK_ADDITIONAL_PATH, sep=";")
bank_additional_full_df = pd.read_csv(BANK_ADDITIONAL_FULL_PATH, sep=";")
```

```
[4]: bank_df
```

```
[4]:
```

	age	job	marital	education	default	balance	housing	loan	\
0	30	unemployed	married	primary	no	1787	no	no	
1	33	services	married	secondary	no	4789	yes	yes	
2	35	management	single	tertiary	no	1350	yes	no	
3	30	management	married	tertiary	no	1476	yes	yes	
4	59	blue-collar	married	secondary	no	0	yes	no	
...	
4516	33	services	married	secondary	no	-333	yes	no	
4517	57	self-employed	married	tertiary	yes	-3313	yes	yes	
4518	57	technician	married	secondary	no	295	no	no	

4519	28	blue-collar	married	secondary	no	1137	no	no
4520	44	entrepreneur	single	tertiary	no	1136	yes	yes

	contact	day	month	duration	campaign	pdays	previous	poutcome	y
0	cellular	19	oct	79	1	-1	0	unknown	no
1	cellular	11	may	220	1	339	4	failure	no
2	cellular	16	apr	185	1	330	1	failure	no
3	unknown	3	jun	199	4	-1	0	unknown	no
4	unknown	5	may	226	1	-1	0	unknown	no
...
4516	cellular	30	jul	329	5	-1	0	unknown	no
4517	unknown	9	may	153	1	-1	0	unknown	no
4518	cellular	19	aug	151	11	-1	0	unknown	no
4519	cellular	6	feb	129	4	211	3	other	no
4520	cellular	3	apr	345	2	249	7	other	no

[4521 rows x 17 columns]

```
[5]: bank_full_df
```

```
[5]:
```

	age	job	marital	education	default	balance	housing	loan	\
0	58	management	married	tertiary	no	2143	yes	no	
1	44	technician	single	secondary	no	29	yes	no	
2	33	entrepreneur	married	secondary	no	2	yes	yes	
3	47	blue-collar	married	unknown	no	1506	yes	no	
4	33	unknown	single	unknown	no	1	no	no	
...
45206	51	technician	married	tertiary	no	825	no	no	
45207	71	retired	divorced	primary	no	1729	no	no	
45208	72	retired	married	secondary	no	5715	no	no	
45209	57	blue-collar	married	secondary	no	668	no	no	
45210	37	entrepreneur	married	secondary	no	2971	no	no	

	contact	day	month	duration	campaign	pdays	previous	poutcome	y
0	unknown	5	may	261	1	-1	0	unknown	no
1	unknown	5	may	151	1	-1	0	unknown	no
2	unknown	5	may	76	1	-1	0	unknown	no
3	unknown	5	may	92	1	-1	0	unknown	no
4	unknown	5	may	198	1	-1	0	unknown	no
...
45206	cellular	17	nov	977	3	-1	0	unknown	yes
45207	cellular	17	nov	456	2	-1	0	unknown	yes
45208	cellular	17	nov	1127	5	184	3	success	yes
45209	telephone	17	nov	508	4	-1	0	unknown	no
45210	cellular	17	nov	361	2	188	11	other	no

[45211 rows x 17 columns]

```
[6]: bank_additional_df
```

```
[6]:      age      job marital      education default housing  loan \
0      30  blue-collar married      basic.9y      no      yes      no
1      39    services  single      high.school      no      no      no
2      25    services married      high.school      no      yes      no
3      38    services married      basic.9y      no  unknown  unknown
4      47      admin. married  university.degree      no      yes      no
...  ...
4114    30      admin. married      basic.6y      no      yes      yes
4115    39      admin. married      high.school      no      yes      no
4116    27    student  single      high.school      no      no      no
4117    58      admin. married      high.school      no      no      no
4118    34  management  single      high.school      no      yes      no
```

```
      contact month day_of_week ... campaign pdays previous \
0      cellular may      fri ...      2      999      0
1      telephone may      fri ...      4      999      0
2      telephone jun      wed ...      1      999      0
3      telephone jun      fri ...      3      999      0
4      cellular nov      mon ...      1      999      0
...  ...
4114    cellular jul      thu ...      1      999      0
4115    telephone jul      fri ...      1      999      0
4116    cellular may      mon ...      2      999      1
4117    cellular aug      fri ...      1      999      0
4118    cellular nov      wed ...      1      999      0
```

```
      poutcome emp.var.rate cons.price.idx cons.conf.idx euribor3m \
0      nonexistent      -1.8      92.893      -46.2      1.313
1      nonexistent      1.1      93.994      -36.4      4.855
2      nonexistent      1.4      94.465      -41.8      4.962
3      nonexistent      1.4      94.465      -41.8      4.959
4      nonexistent      -0.1      93.200      -42.0      4.191
...  ...
4114    nonexistent      1.4      93.918      -42.7      4.958
4115    nonexistent      1.4      93.918      -42.7      4.959
4116      failure      -1.8      92.893      -46.2      1.354
4117    nonexistent      1.4      93.444      -36.1      4.966
4118    nonexistent      -0.1      93.200      -42.0      4.120
```

```
      nr.employed  y
0      5099.1  no
1      5191.0  no
2      5228.1  no
3      5228.1  no
4      5195.8  no
```

```
... ..
4114      5228.1 no
4115      5228.1 no
4116      5099.1 no
4117      5228.1 no
4118      5195.8 no
```

[4119 rows x 21 columns]

```
[7]: bank_additional_full_df
```

```
[7]:      age      job marital      education default housing loan \
0      56  housemaid married      basic.4y      no      no  no
1      57  services married      high.school unknown      no  no
2      37  services married      high.school      no  yes  no
3      40   admin. married      basic.6y      no      no  no
4      56  services married      high.school      no      no  yes
... ..
41183  73   retired married professional.course      no  yes  no
41184  46 blue-collar married professional.course      no      no  no
41185  56   retired married university.degree      no  yes  no
41186  44 technician married professional.course      no      no  no
41187  74   retired married professional.course      no  yes  no
```

```
      contact month day_of_week ... campaign pdays previous \
0  telephone may      mon ...      1  999      0
1  telephone may      mon ...      1  999      0
2  telephone may      mon ...      1  999      0
3  telephone may      mon ...      1  999      0
4  telephone may      mon ...      1  999      0
... ..
41183  cellular nov      fri ...      1  999      0
41184  cellular nov      fri ...      1  999      0
41185  cellular nov      fri ...      2  999      0
41186  cellular nov      fri ...      1  999      0
41187  cellular nov      fri ...      3  999      1
```

```
      poutcome emp.var.rate cons.price.idx cons.conf.idx euribor3m \
0  nonexistent      1.1      93.994      -36.4      4.857
1  nonexistent      1.1      93.994      -36.4      4.857
2  nonexistent      1.1      93.994      -36.4      4.857
3  nonexistent      1.1      93.994      -36.4      4.857
4  nonexistent      1.1      93.994      -36.4      4.857
... ..
41183  nonexistent      -1.1      94.767      -50.8      1.028
41184  nonexistent      -1.1      94.767      -50.8      1.028
41185  nonexistent      -1.1      94.767      -50.8      1.028
```

41186	nonexistent	-1.1	94.767	-50.8	1.028
41187	failure	-1.1	94.767	-50.8	1.028

	nr.employed	y
0	5191.0	no
1	5191.0	no
2	5191.0	no
3	5191.0	no
4	5191.0	no
...
41183	4963.6	yes
41184	4963.6	no
41185	4963.6	no
41186	4963.6	yes
41187	4963.6	no

[41188 rows x 21 columns]

```
[8]: df_list = [bank_df, bank_full_df, bank_additional_df, bank_additional_full_df]
print(f"The bank-full.csv dataset has the shape of: {bank_full_df.shape}")
print(f"It contains the columns: {list(bank_full_df.columns)}\n")
print("-----")
print(f"The bank.csv has the shape of: {bank_df.shape}")
print(f"It contains the columns: {list(bank_df.columns)}\n")
print("-----")
print(f"The bank-additional-full.csv dataset has the shape of: {bank_additional_full_df.shape}")
print(f"It contains the columns: {list(bank_additional_full_df.columns)}\n")
print("-----")
print(f"The bank-additional.csv dataset has the shape of: {bank_additional_df.shape}")
print(f"It contains the columns: {list(bank_additional_df.columns)}\n")
```

The bank-full.csv dataset has the shape of: (45211, 17)

It contains the columns: ['age', 'job', 'marital', 'education', 'default', 'balance', 'housing', 'loan', 'contact', 'day', 'month', 'duration', 'campaign', 'pdays', 'previous', 'poutcome', 'y']

The bank.csv has the shape of: (4521, 17)

It contains the columns: ['age', 'job', 'marital', 'education', 'default', 'balance', 'housing', 'loan', 'contact', 'day', 'month', 'duration', 'campaign', 'pdays', 'previous', 'poutcome', 'y']

The bank-additional-full.csv dataset has the shape of: (41188, 21)

It contains the columns: ['age', 'job', 'marital', 'education', 'default', 'housing', 'loan', 'contact', 'month', 'day_of_week', 'duration', 'campaign',

```
'pdays', 'previous', 'poutcome', 'emp.var.rate', 'cons.price.idx',
'cons.conf.idx', 'euribor3m', 'nr.employed', 'y']
```

```
-----
The bank-additional.csv dataset has the shape of: (4119, 21)
It contains the columns: ['age', 'job', 'marital', 'education', 'default',
'housing', 'loan', 'contact', 'month', 'day_of_week', 'duration', 'campaign',
'pdays', 'previous', 'poutcome', 'emp.var.rate', 'cons.price.idx',
'cons.conf.idx', 'euribor3m', 'nr.employed', 'y']
```

Verification Conclusion: The shapes match the documentation, and the smaller files are confirmed to be true subsets of the larger files. We can now confidently proceed with bank-additional-full.csv for our analysis.

```
[18]: # get full statistical metrics on numerical columns
bank_additional_full_df.describe()
```

```
[18]:
```

	age	duration	campaign	pdays	previous \
count	41188.00000	41188.000000	41188.000000	41188.000000	41188.000000
mean	40.02406	258.285010	2.567593	962.475454	0.172963
std	10.42125	259.279249	2.770014	186.910907	0.494901
min	17.00000	0.000000	1.000000	0.000000	0.000000
25%	32.00000	102.000000	1.000000	999.000000	0.000000
50%	38.00000	180.000000	2.000000	999.000000	0.000000
75%	47.00000	319.000000	3.000000	999.000000	0.000000
max	98.00000	4918.000000	56.000000	999.000000	7.000000

	emp.var.rate	cons.price.idx	cons.conf.idx	euribor3m	nr.employed
count	41188.000000	41188.000000	41188.000000	41188.000000	41188.000000
mean	0.081886	93.575664	-40.502600	3.621291	5167.035911
std	1.570960	0.578840	4.628198	1.734447	72.251528
min	-3.400000	92.201000	-50.800000	0.634000	4963.600000
25%	-1.800000	93.075000	-42.700000	1.344000	5099.100000
50%	1.100000	93.749000	-41.800000	4.857000	5191.000000
75%	1.400000	93.994000	-36.400000	4.961000	5228.100000
max	1.400000	94.767000	-26.900000	5.045000	5228.100000

1.3 Data Cleaning

We'll check for any missing values and duplicates. For that we will need to see what column entries exists for each column in the first place.

```
[9]: # Loop through each column in the dataframe
for column in bank_additional_full_df.columns:
    num_unique_values = bank_additional_full_df[column].nunique()

    print(f"\n----- Column: '{column}' -----")
```

```

print(f"Number of unique values: {num_unique_values}")

# Set a threshold to decide whether to print all unique values
# This avoids printing thousands of unique values for continuous columns
↳like 'age' or 'duration'
if num_unique_values < 15:
    # Sort the values to make them easier to read
    unique_values = sorted(bank_additional_full_df[column].unique())
    print(f"Unique values: {unique_values}")
else:
    # For columns with many unique values, we just note that it's a
    ↳high-cardinality feature
    # We can show a small sample of the unique values
    sample_unique_values = list(bank_additional_full_df[column].unique())[:
    ↳5]
    print(f"Values: [High Cardinality Feature - Sample:
    ↳{sample_unique_values}...]")

```

----- Column: 'age' -----

Number of unique values: 78

Values: [High Cardinality Feature - Sample: [np.int64(56), np.int64(57), np.int64(37), np.int64(40), np.int64(45)]...]

----- Column: 'job' -----

Number of unique values: 12

Unique values: ['admin.', 'blue-collar', 'entrepreneur', 'housemaid', 'management', 'retired', 'self-employed', 'services', 'student', 'technician', 'unemployed', 'unknown']

----- Column: 'marital' -----

Number of unique values: 4

Unique values: ['divorced', 'married', 'single', 'unknown']

----- Column: 'education' -----

Number of unique values: 8

Unique values: ['basic.4y', 'basic.6y', 'basic.9y', 'high.school', 'illiterate', 'professional.course', 'university.degree', 'unknown']

----- Column: 'default' -----

Number of unique values: 3

Unique values: ['no', 'unknown', 'yes']

----- Column: 'housing' -----

Number of unique values: 3

Unique values: ['no', 'unknown', 'yes']

```

----- Column: 'loan' -----
Number of unique values: 3
Unique values: ['no', 'unknown', 'yes']

----- Column: 'contact' -----
Number of unique values: 2
Unique values: ['cellular', 'telephone']

----- Column: 'month' -----
Number of unique values: 10
Unique values: ['apr', 'aug', 'dec', 'jul', 'jun', 'mar', 'may', 'nov', 'oct',
'sep']

----- Column: 'day_of_week' -----
Number of unique values: 5
Unique values: ['fri', 'mon', 'thu', 'tue', 'wed']

----- Column: 'duration' -----
Number of unique values: 1544
Values: [High Cardinality Feature - Sample: [np.int64(261), np.int64(149),
np.int64(226), np.int64(151), np.int64(307)]...]

----- Column: 'campaign' -----
Number of unique values: 42
Values: [High Cardinality Feature - Sample: [np.int64(1), np.int64(2),
np.int64(3), np.int64(4), np.int64(5)]...]

----- Column: 'pdays' -----
Number of unique values: 27
Values: [High Cardinality Feature - Sample: [np.int64(999), np.int64(6),
np.int64(4), np.int64(3), np.int64(5)]...]

----- Column: 'previous' -----
Number of unique values: 8
Unique values: [np.int64(0), np.int64(1), np.int64(2), np.int64(3), np.int64(4),
np.int64(5), np.int64(6), np.int64(7)]

----- Column: 'poutcome' -----
Number of unique values: 3
Unique values: ['failure', 'nonexistent', 'success']

----- Column: 'emp.var.rate' -----
Number of unique values: 10
Unique values: [np.float64(-3.4), np.float64(-3.0), np.float64(-2.9),
np.float64(-1.8), np.float64(-1.7), np.float64(-1.1), np.float64(-0.2),
np.float64(-0.1), np.float64(1.1), np.float64(1.4)]

----- Column: 'cons.price.idx' -----

```



```

Number of unique values: 26
Values: [High Cardinality Feature - Sample: [np.float64(93.994),
np.float64(94.465), np.float64(93.918), np.float64(93.444),
np.float64(93.798)]...]

----- Column: 'cons.conf.idx' -----
Number of unique values: 26
Values: [High Cardinality Feature - Sample: [np.float64(-36.4),
np.float64(-41.8), np.float64(-42.7), np.float64(-36.1), np.float64(-40.4)]...]

----- Column: 'euribor3m' -----
Number of unique values: 316
Values: [High Cardinality Feature - Sample: [np.float64(4.857),
np.float64(4.856), np.float64(4.855), np.float64(4.859), np.float64(4.86)]...]

----- Column: 'nr.employed' -----
Number of unique values: 11
Unique values: [np.float64(4963.6), np.float64(4991.6), np.float64(5008.7),
np.float64(5017.5), np.float64(5023.5), np.float64(5076.2), np.float64(5099.1),
np.float64(5176.3), np.float64(5191.0), np.float64(5195.8), np.float64(5228.1)]

----- Column: 'y' -----
Number of unique values: 2
Unique values: ['no', 'yes']

```

1.4 1.4. Feature Analysis: Feature vs. Target

Now we'll analyze how each feature relates to the subscription outcome *y*. This will help us identify potentially predictive features.

1.4.1 1.4.1 Categorical Features vs. Target ('y')

```

[10]: # List of key categorical features to analyze
cat_features_to_plot = ['job', 'marital', 'education', 'default', 'housing',
    ↪ 'loan', 'contact', 'poutcome']

# Create plots
fig, axes = plt.subplots(4, 2, figsize=(20, 25))
axes = axes.flatten()


for i, col in enumerate(cat_features_to_plot):
    # Use hue to show the distribution of the target variable for each category
    sns.countplot(y=col, data=bank_additional_full_df, ax=axes[i],
    ↪ order=bank_additional_full_df[col].value_counts().index, hue='y',
    ↪ palette='viridis')
    axes[i].set_title(f"{col.capitalize()} vs. Subscription", fontsize=14)
    axes[i].set_xlabel('Count')
    axes[i].set_ylabel('')

```

```
axes[i].legend(title='Subscribed')
```

```
plt.tight_layout()
```

```
plt.show()
```



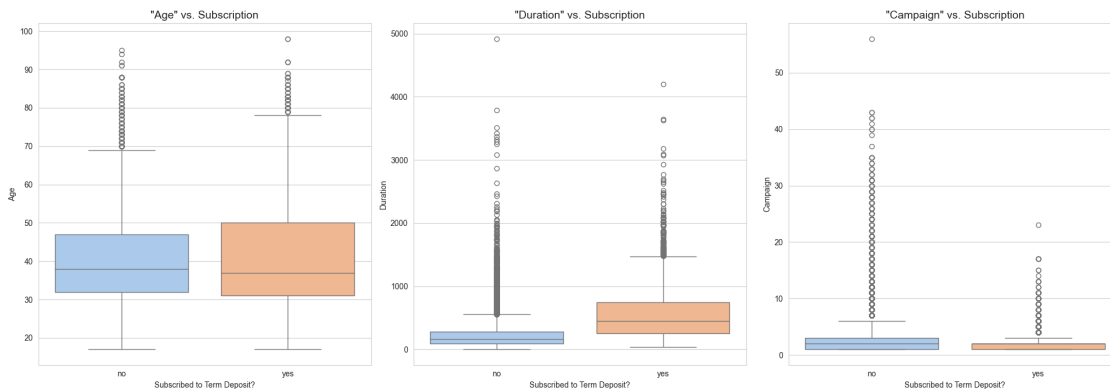
1.4.2 1.4.2 Numerical Features vs. Target ('y')

```
[11]: # List of key numerical features to analyze
num_features_to_plot = ['age', 'duration', 'campaign']

# Create boxplots
fig, axes = plt.subplots(1, 3, figsize=(20, 7))

for i, col in enumerate(num_features_to_plot):
    sns.boxplot(x='y', y=col, data=bank_additional_full_df, ax=axes[i], hue="y",
                palette='pastel', legend=False)
    axes[i].set_title(f'"{col.capitalize()}" vs. Subscription', fontsize=14)
    axes[i].set_xlabel('Subscribed to Term Deposit?')
    axes[i].set_ylabel(col.capitalize())

plt.tight_layout()
plt.show()
```



1.4.3 1.4.3 Time & Previous Campaign Features vs. Target ('y')

```
[12]: # Features to analyze
features_to_plot = ['month', 'day_of_week', 'previous']

# Define a chronological order for months and days
month_order = ['mar', 'apr', 'may', 'jun', 'jul', 'aug', 'sep', 'oct', 'nov', 'dec']
day_order = ['mon', 'tue', 'wed', 'thu', 'fri']
order_map = {'month': month_order, 'day_of_week': day_order, 'previous': sorted(bank_additional_full_df['previous'].unique())}


fig, axes = plt.subplots(1, 3, figsize=(22, 7))
axes = axes.flatten()
```

```

for i, col in enumerate(features_to_plot):
    sns.countplot(x=col, data=bank_additional_full_df, ax=axes[i],
        order=order_map[col], hue='y', palette='viridis')
    axes[i].set_title(f"{col.capitalize()} vs. Subscription", fontsize=14)
    axes[i].set_ylabel('Count')
    axes[i].set_xlabel('')
    axes[i].tick_params(axis='x', rotation=45) # Rotate labels for readability
    axes[i].legend(title='Subscribed')

plt.tight_layout()
plt.show()

```



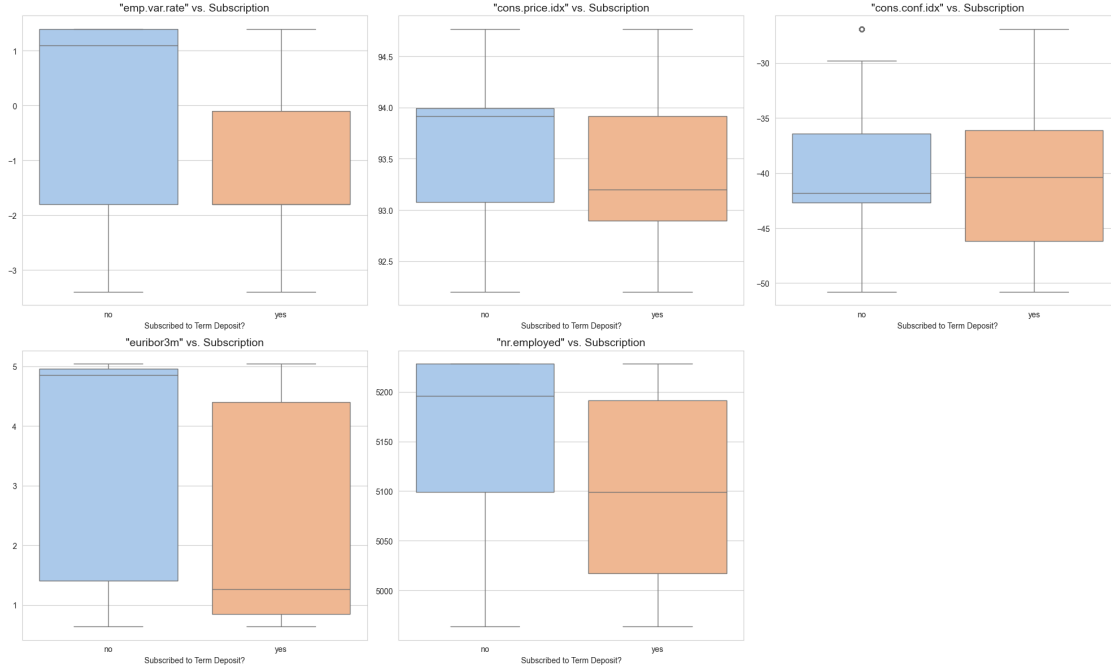
1.4.4 Social & Economic Features vs. Target ('y')

```

[13]: # List of social and economic features
social_econ_features = ['emp.var.rate', 'cons.price.idx', 'cons.conf.idx',
    'euribor3m', 'nr.employed']

# Create boxplots
fig, axes = plt.subplots(2, 3, figsize=(20, 12))
axes = axes.flatten()
for i, col in enumerate(social_econ_features):
    sns.boxplot(x='y', y=col, data=bank_additional_full_df, ax=axes[i],
        palette='pastel', hue='y', legend=False)
    axes[i].set_title(f"{col} vs. Subscription", fontsize=14)
    axes[i].set_xlabel('Subscribed to Term Deposit?')
    axes[i].set_ylabel('')
# Hide the empty subplot
fig.delaxes(axes[5])
plt.tight_layout()
plt.show()

```



1.4.5 1.4.5 Visual Feature Analysis Summary

A. Categorical Features vs. Target ('y')

- **Job & Education:** Administrative staff, technicians, and blue-collar workers form the largest groups of clients contacted. However, students and retired individuals show a proportionally higher subscription rate ("yes") compared to other job categories. Clients with a university degree were contacted most frequently, but subscription rates appear relatively consistent across different education levels, with a slight increase for those with higher education.
- **Marital & Default Status:** Married individuals are the largest client segment, followed by singles. The "Default" status is overwhelmingly "no," and very few clients with a "yes" status subscribed. This suggests that clients with a history of credit default are unlikely to subscribe.
- **Housing & Loan:** The subscription rate is higher for clients who do not have an existing housing loan. Similarly, clients without a personal loan are more likely to subscribe than those with one.
- **Contact & Previous Campaign Outcome (Poutcome):** Contacting clients via "cellular" is associated with a much higher subscription rate than "telephone". As expected, a "success" outcome from a previous campaign is a strong predictor of a "yes" for the current campaign.

B. Numerical Features vs. Target ('y')

- **Age:** The age distribution for both subscribers and non-subscribers is similar, with the median age for subscribers appearing slightly higher than for non-subscribers.
- **Duration:** The duration of the last contact is significantly higher for clients who subscribed. This is a key insight but must be handled with care, as call duration is not known until after the call is made. Thus, it cannot be used as a predictive feature for a pre-call model.

- **Campaign:** The number of contacts during the campaign is heavily skewed towards the lower end for both groups. However, the median number of contacts for subscribers is slightly lower than for non-subscribers, suggesting that fewer contacts are often more effective.

C. Time & Previous Campaign Features vs. Target ('y')

- **Month & Day of the Week:** Subscription rates vary significantly by month, with the highest success rates appearing in March, September, October, and December. The day of the week does not show a significant variation in subscription rates.
- **Previous:** A higher number of previous contacts (before the current campaign) is associated with a higher likelihood of subscribing, although the vast majority of clients had no previous contact.

D. Social & Economic Features vs. Target ('y')

- The boxplots for social and economic indicators show clear distinctions between subscribers and non-subscribers:
 - **Subscribers (“yes”) are associated with:**
 - * Lower (more negative) employment variation rates (`emp.var.rate`).
 - * Lower consumer price indexes (`cons.price.idx`).
 - * Higher (less negative) consumer confidence indexes (`cons.conf.idx`).
 - * Lower 3-month Euribor rates (`euribor3m`).
 - * Lower numbers of employees (`nr.employed`).
- These trends suggest that clients are more likely to subscribe during periods of lower economic pressure (e.g., lower interest rates, lower employment figures, and higher consumer confidence).

1.5 1.5 Quantifying Feature Relationships

The next logical step in the Exploratory Data Analysis (EDA) would be to quantify the relationships observed visually in the last section. The plots have shown a good intuition about the data, and the next step is to generate concrete numbers to support these findings before moving to preprocessing the data.

1.5.1 1.5.1 Correlation of numerical Columns


On how to interpret this correlation matrix refer to this [website](#) which explains how to read a correlation matrix.

```
[16]: # Select only numerical columns for correlation
numerical_cols = bank_additional_full_df.select_dtypes(include=np.number).
      ↪ columns

# Calculate the correlation matrix
corr_matrix = bank_additional_full_df[numerical_cols].corr()

# Plot the heatmap
plt.figure(figsize=(12, 10))
sns.heatmap(corr_matrix, annot=True, cmap='coolwarm', fmt=".2f")
plt.title('Correlation Matrix of Numerical Features')
```

```
plt.show()
```



1.5.2 Subscription Rate Analysis

```
[21]: # First, create a copy and convert 'y' to a numerical format 1: yes 0: no
df_rate = bank_additional_full_df.copy()
df_rate['y_numeric'] = df_rate['y'].apply(lambda x: 1 if x == 'yes' else 0)

[22]: # We will reuse the df_rate DataFrame with the 'y_numeric' column
cat_features = ['job', 'marital', 'education', 'contact', 'poutcome', 'month']

fig, axes = plt.subplots(3, 2, figsize=(20, 20))
axes = axes.flatten()
fig.suptitle('Subscription Rate by Categorical Features', fontsize=16)
```

```

for i, col in enumerate(cat_features):
    # Calculate subscription rate
    rate = df_rate.groupby(col)['y_numeric'].mean().sort_values(ascending=True)

    # Plot
    rate.plot(kind='barh', ax=axes[i], color=sns.color_palette('viridis',
↪len(rate)))
    axes[i].set_title(f'Subscription Rate by {col.capitalize()}')
    axes[i].set_xlabel('Subscription Rate (%)')
    axes[i].set_ylabel('')
    # Format x-axis as percentage
    axes[i].xaxis.set_major_formatter(plt.FuncFormatter(lambda x, _: '{:.0%}'.
↪format(x)))

plt.tight_layout(rect=[0, 0, 1, 0.96])
plt.show()

```


Subscription Rate by Categorical Features



1.5.3 Correlation Analysis Summary

Correlation Matrix of Numerical Features The correlation matrix reveals strong relationships between several of the socio-economic indicator variables. - **High Multicollinearity:** There are very strong positive correlations between `emp.var.rate`, `euribor3m`, and `nr.employed`, with correlation coefficients ranging from 0.91 to 0.97. This indicates that these features measure similar underlying economic conditions. For certain models, like logistic regression, this multicollinearity can be problematic, and we may consider removing some of these redundant features during the feature selection phase. - **Other Correlations:** A moderate negative correlation exists between `previous` (number of contacts before this campaign) and `pdays` (days since last contact), with a coefficient of -0.59. This makes sense, as clients with more previous contacts are likely to have been contacted more recently.

Subscription Rate Analysis Calculating the exact subscription rates confirms the visual insights from the earlier plots and provides precise metrics.

- By Job: Students (31.4%) and retired clients (25.2%) have the highest likelihood of subscribing to a term deposit. In contrast, blue-collar workers (6.9%) have the lowest subscription rate. This reinforces that targeting specific client professions could significantly increase campaign efficiency.
- By Previous Outcome: The outcome of a previous campaign is an extremely powerful indicator. Clients with a “success” in a prior campaign have a 65.1% subscription rate, which is dramatically higher than for clients with no previous contact (“nonexistent” at 8.8%) or a previous “failure” (14.2%).

1.6 Conclusion of Exploratory Data Analysis (EDA)

This EDA has identified key patterns and characteristics within the dataset that will be needed during our modeling strategy.

1. Strong Predictors Identified: Several features show a strong relationship with the client’s decision to subscribe. The most influential appear to be the outcome of the previous campaign (`poutcome`), the month of contact, the contact method (`contact`), and the socio-economic indicators (`emp.var.rate`, `euribor3m`, etc.). Client job type also shows significant variance in subscription rates.
2. Data Leakage for Model: The duration feature is highly correlated with the outcome but must be excluded from the predictive model, since this information is not available before an actual call is made.
3. Data Quality & Preprocessing Needs:
 - “Unknown” Values: Several key categorical features contain “unknown” entries, which will need to be handled, either by treating them as a distinct category, imputing them with different categories or replacing as NAN values.
 - Class Imbalance: The target variable ‘y’ is highly imbalanced, with a large majority of non-subscribers. This must be addressed during the modeling phase to prevent model bias.
 - Multicollinearity: The high correlation among socio-economic features suggests that feature selection is an important step.
 - When features are highly correlated, it becomes difficult to distinguish their individual effects on the target variable.
 - Multicollinearity can be problematic for certain models, such as logistic regression
 - Highly correlated features provide redundant information, thus by removing them we can reduce model complexity

2 Data Preprocessing

Following the EDA, the next phase is Data Preprocessing. While the EDA serves as a diagnostic investigation, providing a deep understanding of the dataset’s structure, underlying patterns, and most importantly, its limitations and potential issues.

The analysis revealed mentioned challenges that must be addressed before modeling such as: - the presence of “unknown” values in categorical features and the encoding of categorical features -

significant class imbalance in the target variable - high multicollinearity among the socio-economic indicators - data leakage risk from the duration feature

```
[95]: df = bank_additional_full_df.copy()
df
```

```
[95]:      age      job marital      education default housing loan \
0      56  housemaid married      basic.4y      no      no      no
1      57  services married      high.school unknown      no      no
2      37  services married      high.school      no      yes      no
3      40   admin. married      basic.6y      no      no      no
4      56  services married      high.school      no      no      yes
...  ...
41183  73    retired married professional.course      no      yes      no
41184  46 blue-collar married professional.course      no      no      no
41185  56    retired married university.degree      no      yes      no
41186  44  technician married professional.course      no      no      no
41187  74    retired married professional.course      no      yes      no
```

```
      contact month day_of_week ... campaign pdays previous \
0  telephone may      mon ...      1  999      0
1  telephone may      mon ...      1  999      0
2  telephone may      mon ...      1  999      0
3  telephone may      mon ...      1  999      0
4  telephone may      mon ...      1  999      0
...  ...
41183  cellular nov      fri ...      1  999      0
41184  cellular nov      fri ...      1  999      0
41185  cellular nov      fri ...      2  999      0
41186  cellular nov      fri ...      1  999      0
41187  cellular nov      fri ...      3  999      1
```

```
      poutcome emp.var.rate cons.price.idx cons.conf.idx euribor3m \
0  nonexistent      1.1      93.994      -36.4      4.857
1  nonexistent      1.1      93.994      -36.4      4.857
2  nonexistent      1.1      93.994      -36.4      4.857
3  nonexistent      1.1      93.994      -36.4      4.857
4  nonexistent      1.1      93.994      -36.4      4.857
...  ...
41183  nonexistent      -1.1      94.767      -50.8      1.028
41184  nonexistent      -1.1      94.767      -50.8      1.028
41185  nonexistent      -1.1      94.767      -50.8      1.028
41186  nonexistent      -1.1      94.767      -50.8      1.028
41187    failure      -1.1      94.767      -50.8      1.028
```

```
      nr.employed y
0      5191.0 no
```

```

1          5191.0  no
2          5191.0  no
3          5191.0  no
4          5191.0  no
...
41183      4963.6  yes
41184      4963.6  no
41185      4963.6  no
41186      4963.6  yes
41187      4963.6  no

```

```
[41188 rows x 21 columns]
```

2.1 2.1 Address Data Leakage by Removing ‘duration’ column

```
[96]: # Remove the 'duration' column to prevent data leakage
df.drop('duration', axis=1, inplace=True)
```

2.2 2.2 Encode the Target Variable ‘y’

Machine learning models require numerical inputs, so our first and simplest step is to convert the y column’s values from “yes” and “no” to a binary format (1 and 0).

```
[97]: df['y'] = df['y'].apply(lambda x: 1 if x == 'yes' else 0)
df
```

```
[97]:
```

	age	job	marital	education	default	housing	loan	\
0	56	housemaid	married	basic.4y	no	no	no	
1	57	services	married	high.school	unknown	no	no	
2	37	services	married	high.school	no	yes	no	
3	40	admin.	married	basic.6y	no	no	no	
4	56	services	married	high.school	no	no	yes	
...	
41183	73	retired	married	professional.course	no	yes	no	
41184	46	blue-collar	married	professional.course	no	no	no	
41185	56	retired	married	university.degree	no	yes	no	
41186	44	technician	married	professional.course	no	no	no	
41187	74	retired	married	professional.course	no	yes	no	

	contact	month	day_of_week	campaign	pdays	previous	poutcome	\
0	telephone	may	mon	1	999	0	nonexistent	
1	telephone	may	mon	1	999	0	nonexistent	
2	telephone	may	mon	1	999	0	nonexistent	
3	telephone	may	mon	1	999	0	nonexistent	
4	telephone	may	mon	1	999	0	nonexistent	
...	
41183	cellular	nov	fri	1	999	0	nonexistent	

41184	cellular	nov	fri	1	999	0	nonexistent
41185	cellular	nov	fri	2	999	0	nonexistent
41186	cellular	nov	fri	1	999	0	nonexistent
41187	cellular	nov	fri	3	999	1	failure

	emp.var.rate	cons.price.idx	cons.conf.idx	euribor3m	nr.employed	y
0	1.1	93.994	-36.4	4.857	5191.0	0
1	1.1	93.994	-36.4	4.857	5191.0	0
2	1.1	93.994	-36.4	4.857	5191.0	0
3	1.1	93.994	-36.4	4.857	5191.0	0
4	1.1	93.994	-36.4	4.857	5191.0	0
...
41183	-1.1	94.767	-50.8	1.028	4963.6	1
41184	-1.1	94.767	-50.8	1.028	4963.6	0
41185	-1.1	94.767	-50.8	1.028	4963.6	0
41186	-1.1	94.767	-50.8	1.028	4963.6	1
41187	-1.1	94.767	-50.8	1.028	4963.6	0

[41188 rows x 20 columns]

Why do we need to encode categoriacal columns? Machine learning algorithms are based on mathematical equations and can only process numerical data. They cannot understand text labels like 'yes' or 'no' directly. In this case we the y category with 1 and 0 since it has only 2 options. **But what about features that have more categories?**

2.3 Encoding categoriacal features with more than 2 categories using One-Hot-Encoding

The principle of one-hot encoding works exactly the same way, whether a feature has two categories or many more. It simply expands to create a new binary column for every unique category.

So let's encode all categorical columns.

```
[98]: # 1. Identify all categorical columns that need encoding
# These are the columns with string values (object dtype)
categorical_features = df.select_dtypes(include=['object']).columns
categorical_features
```

```
[98]: Index(['job', 'marital', 'education', 'default', 'housing', 'loan', 'contact',
          'month', 'day_of_week', 'poutcome'],
          dtype='object')
```

```
[99]: # 2. Apply one-hot encoding using pd.get_dummies()
df_encoded = pd.get_dummies(df, columns=categorical_features,
                             drop_first=False, dtype=int)

# 3. Display the results
```

```
print(f"Original shape of the DataFrame: {df.shape}")
print(f"Shape after one-hot encoding: {df_encoded.shape}\n")
```

Original shape of the DataFrame: (41188, 20)

Shape after one-hot encoding: (41188, 63)

```
[100]: df_encoded.head(5)
```

```
[100]:
```

	age	campaign	pdays	previous	emp.var.rate	cons.price.idx	\
0	56	1	999	0	1.1	93.994	
1	57	1	999	0	1.1	93.994	
2	37	1	999	0	1.1	93.994	
3	40	1	999	0	1.1	93.994	
4	56	1	999	0	1.1	93.994	

	cons.conf.idx	euribor3m	nr.employed	y	...	month_oct	month_sep	\
0	-36.4	4.857	5191.0	0	...	0	0	
1	-36.4	4.857	5191.0	0	...	0	0	
2	-36.4	4.857	5191.0	0	...	0	0	
3	-36.4	4.857	5191.0	0	...	0	0	
4	-36.4	4.857	5191.0	0	...	0	0	

	day_of_week_fri	day_of_week_mon	day_of_week_thu	day_of_week_tue	\
0	0	1	0	0	
1	0	1	0	0	
2	0	1	0	0	
3	0	1	0	0	
4	0	1	0	0	

	day_of_week_wed	poutcome_failure	poutcome_nonexistent	poutcome_success
0	0	0	1	0
1	0	0	1	0
2	0	0	1	0
3	0	0	1	0
4	0	0	1	0

[5 rows x 63 columns]

2.4 2.4 Separating Feature and Target Variable

```
[101]: # Separate the features (X) from the target variable (y)
X = df_encoded.drop('y', axis=1)
y = df_encoded['y']

print(f"Shape of features (X): {X.shape}")
print(f"Shape of target (y): {y.shape}")
```

Shape of features (X): (41188, 62)
Shape of target (y): (41188,)

2.5 Split Data into Training and Testing Sets

```
[102]: from sklearn.model_selection import train_test_split

# Split the data into training and testing sets (80% train, 20% test)
# 'stratify=y' ensures that the proportion of subscribers is the same in both
    ↪ sets
X_train, X_test, y_train, y_test = train_test_split(X, y, test_size=0.2,
    ↪ random_state=42, stratify=y)

print(f"Training set size: {X_train.shape[0]} samples")
print(f"Testing set size: {X_test.shape[0]} samples")
```

Training set size: 32950 samples
Testing set size: 8238 samples

3. Creating a Baseline Model using LogisticRegression for this Classification Task

```
[103]: from sklearn.linear_model import LogisticRegression
# 1. Initialize the Logistic Regression model
# We'll set max_iter to a higher value to ensure the model's algorithm
    ↪ converges.
log_reg_baseline = LogisticRegression(random_state=42, max_iter=500)
```

```
[104]: # 2. Train the model on the (unscaled, imbalanced) training data
print("Training the baseline logistic regression model...")
log_reg_baseline.fit(X_train, y_train)
print("Model training complete.")
```

Training the baseline logistic regression model...

/Users/tomyle/Pycharm/xai_capstone/.venv/lib/python3.12/site-packages/sklearn/linear_model/_linear_loss.py:200: RuntimeWarning: divide by zero encountered in matmul

```
raw_prediction = X @ weights + intercept
```

/Users/tomyle/Pycharm/xai_capstone/.venv/lib/python3.12/site-packages/sklearn/linear_model/_linear_loss.py:200: RuntimeWarning: overflow encountered in matmul

```
raw_prediction = X @ weights + intercept
```

/Users/tomyle/Pycharm/xai_capstone/.venv/lib/python3.12/site-packages/sklearn/linear_model/_linear_loss.py:200: RuntimeWarning: invalid value encountered in matmul

```
raw_prediction = X @ weights + intercept
```

/Users/tomyle/Pycharm/xai_capstone/.venv/lib/python3.12/site-

```

packages/sklearn/linear_model/_linear_loss.py:330: RuntimeWarning: divide by
zero encountered in matmul
    grad[:n_features] = X.T @ grad_pointwise + l2_reg_strength * weights
/Users/tomyle/Pycharm/xai_capstone/.venv/lib/python3.12/site-
packages/sklearn/linear_model/_linear_loss.py:330: RuntimeWarning: overflow
encountered in matmul
    grad[:n_features] = X.T @ grad_pointwise + l2_reg_strength * weights
/Users/tomyle/Pycharm/xai_capstone/.venv/lib/python3.12/site-
packages/sklearn/linear_model/_linear_loss.py:330: RuntimeWarning: invalid value
encountered in matmul
    grad[:n_features] = X.T @ grad_pointwise + l2_reg_strength * weights
Model training complete.

/Users/tomyle/Pycharm/xai_capstone/.venv/lib/python3.12/site-
packages/sklearn/linear_model/_logistic.py:470: ConvergenceWarning: lbfgs failed
to converge after 500 iteration(s) (status=1):
STOP: TOTAL NO. OF ITERATIONS REACHED LIMIT

```

Increase the number of iterations to improve the convergence (max_iter=500).
You might also want to scale the data as shown in:

```

https://scikit-learn.org/stable/modules/preprocessing.html
Please also refer to the documentation for alternative solver options:
https://scikit-learn.org/stable/modules/linear\_model.html#logistic-
regression
    n_iter_i = _check_optimize_result(

```

```

[105]: # 3. Make predictions on the test data
y_pred_baseline = log_reg_baseline.predict(X_test)
print("\nPredictions have been made on the test set.")

```

Predictions have been made on the test set.

```

/Users/tomyle/Pycharm/xai_capstone/.venv/lib/python3.12/site-
packages/sklearn/utils/extmath.py:203: RuntimeWarning: divide by zero
encountered in matmul
    ret = a @ b
/Users/tomyle/Pycharm/xai_capstone/.venv/lib/python3.12/site-
packages/sklearn/utils/extmath.py:203: RuntimeWarning: overflow encountered in
matmul
    ret = a @ b
/Users/tomyle/Pycharm/xai_capstone/.venv/lib/python3.12/site-
packages/sklearn/utils/extmath.py:203: RuntimeWarning: invalid value encountered
in matmul
    ret = a @ b

```


3.1 3. 1 Model Evaluation

3.1.1 3.1.1 Confusion Matrix

```
[106]: from sklearn.metrics import classification_report, confusion_matrix
# --- 1. Confusion Matrix ---
print("--- Confusion Matrix ---")
cm = confusion_matrix(y_test, y_pred_baseline)


# For a nicer plot
plt.figure(figsize=(6, 4))
sns.heatmap(cm, annot=True, fmt='d', cmap='Blues',
            xticklabels=['Predicted No', 'Predicted Yes'],
            yticklabels=['Actual No', 'Actual Yes'])
plt.title('Baseline Model Confusion Matrix')
plt.show()

# --- 2. Classification Report ---
print("\n--- Classification Report ---")
report = classification_report(y_test, y_pred_baseline, target_names=['No',
                               'Yes (Class 1)'])
print(report)

TN, FP, FN, TP = cm.ravel()
# 2. Calculate the metrics using their formulas
# Accuracy: (TP + TN) / Total
accuracy = (TP + TN) / (TP + TN + FP + FN)
# Precision: TP / (TP + FP)
# How many of the predicted positives were actually positive?
precision = TP / (TP + FP)
# Sensitivity (Recall or True Positive Rate): TP / (TP + FN)
# How many of the actual positives were correctly identified?
sensitivity = TP / (TP + FN)
# Specificity (True Negative Rate): TN / (TN + FP)
# How many of the actual negatives were correctly identified?
specificity = TN / (TN + FP)

# 3. Print the results in a clear format
print("--- Detailed Metrics ---")
print(f"Accuracy:    {accuracy:.4f} (Overall correctness)")
print(f"Precision:    {precision:.4f} (Correctness of positive predictions)")
print(f"Sensitivity: {sensitivity:.4f} (Ability to find all positive samples - a.k.a. Recall)")
print(f"Specificity: {specificity:.4f} (Ability to find all negative samples)")

--- Confusion Matrix ---
```



--- Classification Report ---

	precision	recall	f1-score	support
No (Class 0)	0.91	0.99	0.95	7310
Yes (Class 1)	0.71	0.21	0.32	928
accuracy			0.90	8238
macro avg	0.81	0.60	0.64	8238
weighted avg	0.89	0.90	0.88	8238

--- Detailed Metrics ---

Accuracy: 0.9012 (Overall correctness)
Precision: 0.7065 (Correctness of positive predictions)
Sensitivity: 0.2101 (Ability to find all positive samples - a.k.a. Recall)
Specificity: 0.9889 (Ability to find all negative samples)

The Goal is to improve the TP and TN

3.1.2 3.1.2 Roc Curve and ROCAUC

```
[107]: from sklearn.metrics import roc_curve, roc_auc_score
# 1. Get prediction probabilities for the positive class (Class 1)
y_pred_probs = log_reg_baseline.predict_proba(X_test)[: , 1]
```

```

# 2. Calculate the ROC curve points
fpr, tpr, thresholds = roc_curve(y_test, y_pred_probs)

# 3. Calculate the ROC AUC score
roc_auc = roc_auc_score(y_test, y_pred_probs)
print(f"ROC AUC Score: {roc_auc:.4f}")

# 4. Plot the ROC curve
plt.figure(figsize=(8, 6))
plt.plot(fpr, tpr, color='blue', lw=2, label=f'ROC curve (area = {roc_auc:.2f})')
plt.plot([0, 1], [0, 1], color='red', lw=2, linestyle='--', label='No-Skill Line')
plt.xlim([0.0, 1.0])
plt.ylim([0.0, 1.05])
plt.xlabel('False Positive Rate')
plt.ylabel('True Positive Rate')
plt.title('Receiver Operating Characteristic (ROC) Curve')
plt.legend(loc="lower right")
plt.show()

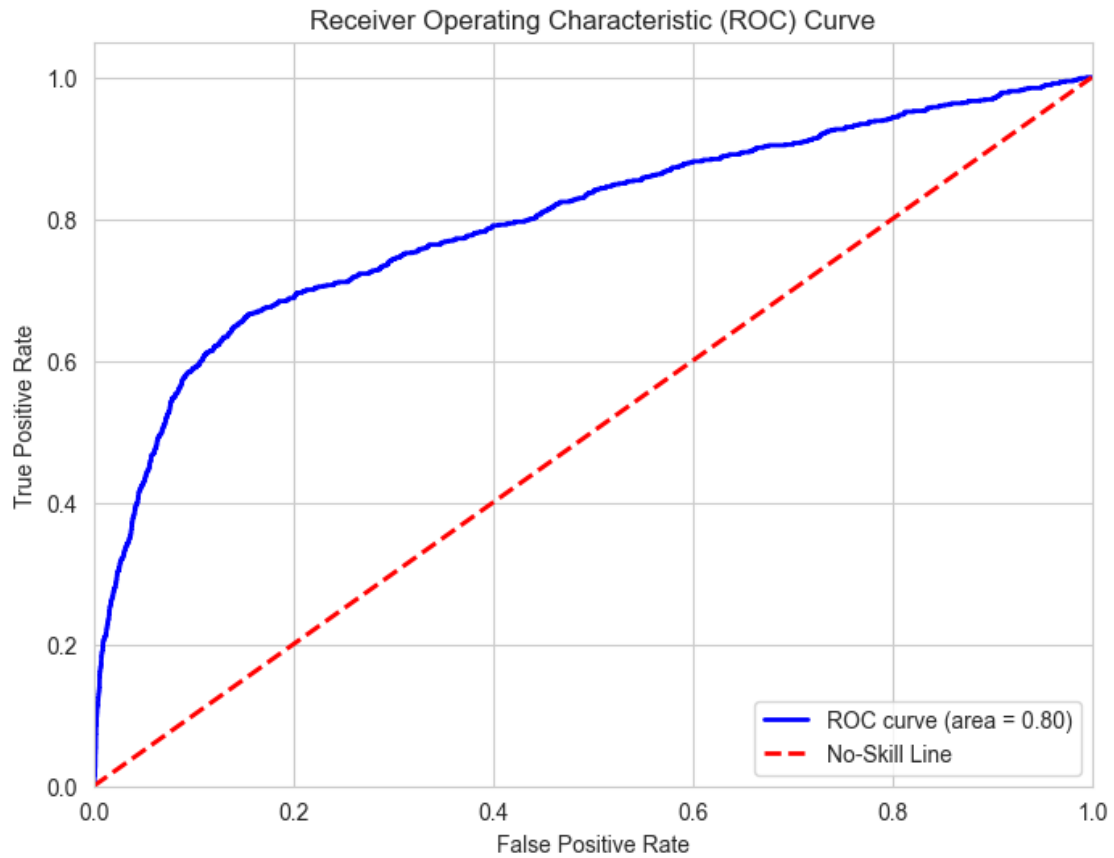
```

ROC AUC Score: 0.7967

```

/Users/tomyle/Pycharm/xai_capstone/.venv/lib/python3.12/site-
packages/sklearn/utils/extmath.py:203: RuntimeWarning: divide by zero
encountered in matmul
  ret = a @ b
/Users/tomyle/Pycharm/xai_capstone/.venv/lib/python3.12/site-
packages/sklearn/utils/extmath.py:203: RuntimeWarning: overflow encountered in
matmul
  ret = a @ b
/Users/tomyle/Pycharm/xai_capstone/.venv/lib/python3.12/site-
packages/sklearn/utils/extmath.py:203: RuntimeWarning: invalid value encountered
in matmul
  ret = a @ b

```



4 Explainable AI (XAI)

4.1 Global Explainability

4.2 Local Explainability

[]:

[]: